

Initial BRD Intake vs Final BRD Output

This table is extremely useful for:

- Explaining **AI value to leadership**
- Defining **Human-in-the-Loop boundaries**
- Training BAs and Product Managers

Dimension	Initial BRD Intake (AI Input)	Final BRD Output (Human Review)
Purpose	Capture raw intent	Act as delivery contract
Owner	Business + AI	Business (human authority)
Language	Free-form, messy, incomplete	Precise, testable, structured
Completeness	Partial allowed	Mandatory completeness
Ambiguity	Expected	Not allowed
Open Questions	Allowed & encouraged	Must be zero or deferred
Requirements	Unstructured bullets	Atomic, ID-based FRs
NFRs	Often missing	Explicit & categorized
Acceptance Criteria	Optional	Mandatory
Traceability	None or partial	100% enforced
Version Control	Informal	Formal
AI behaviour	Ask, infer, normalize	Stop asking, synthesize
Human role	Provide intent & answers	Review, approve, own

Lifecycle View

Initial BRD Intake (DOCX)



AI Agents (AE-Discover)

- Normalize
- Detect gaps
- Ask questions
- Iterate



✓ Final BRD Output (DOCX) ← HUMAN SIGN-OFF



AE-Design

- LLD
- UX
- Architecture

Initial BRD Intake

- Collected **once per initiative**
- Can be updated incrementally
- Lives early in Discover phase
- Accepts chaos

✓ Final BRD Output

- Generated **once convergence rules are met**
- Immutable after sign-off (versioned only)
- Feeds delivery & design
- Zero ambiguity tolerance

5 Why This Design Is Correct (Architecturally)

- Prevents **LLM hallucinations**
- Preserves **business accountability**

- Enables **repeatable AI reasoning**
- Aligns with **regulated enterprise SDLC**
- Scales across teams & programs